

# MLP on MNIST dataset in Keras

## Steps to follow in this assignment:-

1. Load libraries and MNIST data
2. Normalize the data
3. Convert class label values to one-hot encoded values
4. Implementing Softmax classifier with 2 , 3 , 5 hidden layers
5. Implementing model with and without dropout and normalization
6. Compiling the model
7. Plotting Categorical Crossentropy Loss VS Epochs plot
8. Plotting Violin plots to check weights distribution

## 1. Loading libraries

```
In [0]: from keras.utils import np_utils
        from keras.datasets import mnist
        import seaborn as sns
        from keras.initializers import RandomNormal
```

Using TensorFlow backend.

## 2. Loading MNIST data

```
In [0]: # Loading MNIST Train and Test data
        (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 0s 0us/step

```
In [0]: # Printing to see how many data points are there for train and test and to see
        each image size
```

```
print("Number of Train data points :", X_train.shape[0], "and each image is of
shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of Test data points :", X_test.shape[0], "and each image is of s
hape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of Train data points : 60000 and each image is of shape (28, 28)  
Number of Test data points : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
# We are reshaping the train and test data points to 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1] * X_train.shape[2]
])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1] * X_test.shape[2])
```

```
In [0]: # after converting the input images from 3d to 2d vectors , printing the shape
s of train and test data
```

```
print("Number of Train data points after reshape :", X_train.shape[0], "and ea
ch image is of shape (%d)"%(X_train.shape[1]))
print("Number of Test data points after reshape :", X_test.shape[0], "and each
image is of shape (%d)"%(X_test.shape[1]))
```

```
Number of Train data points after reshape : 60000 and each image is of shape
(784)
```

```
Number of Test data points after reshape : 10000 and each image is of shape
(784)
```

In [0]: *# Printing an example data point*

```
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
```

## 2.1 Normalizing the data

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets normalize the data
# so that all values lie between 0 to 1 instead of 0 to 255
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$  . This is min-max type of normalization

X_train = X_train/255
X_test = X_test/255
```

```
In [0]: # printing example train data point after normlizing  
print(X_train[0])
```

|            |            |            |            |            |            |
|------------|------------|------------|------------|------------|------------|
| [0.        | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.01176471 | 0.07058824 | 0.07058824 | 0.07058824 |
| 0.49411765 | 0.53333333 | 0.68627451 | 0.10196078 | 0.65098039 | 1.         |
| 0.96862745 | 0.49803922 | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.11764706 | 0.14117647 | 0.36862745 | 0.60392157 |
| 0.66666667 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 |
| 0.88235294 | 0.6745098  | 0.99215686 | 0.94901961 | 0.76470588 | 0.25098039 |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.19215686 |
| 0.93333333 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 |
| 0.99215686 | 0.99215686 | 0.99215686 | 0.98431373 | 0.36470588 | 0.32156863 |
| 0.32156863 | 0.21960784 | 0.15294118 | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.07058824 | 0.85882353 | 0.99215686 |
| 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 | 0.77647059 | 0.71372549 |
| 0.96862745 | 0.94509804 | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.31372549 | 0.61176471 | 0.41960784 | 0.99215686 |
| 0.99215686 | 0.80392157 | 0.04313725 | 0.         | 0.16862745 | 0.60392157 |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.05490196 | 0.00392157 | 0.60392157 | 0.99215686 | 0.35294118 |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.54509804 | 0.99215686 | 0.74509804 | 0.00784314 | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |

|            |            |            |            |            |            |
|------------|------------|------------|------------|------------|------------|
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.04313725 |
| 0.74509804 | 0.99215686 | 0.2745098  | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.1372549  | 0.94509804 |
| 0.88235294 | 0.62745098 | 0.42352941 | 0.00392157 | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.31764706 | 0.94117647 | 0.99215686 |
| 0.99215686 | 0.46666667 | 0.09803922 | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.17647059 | 0.72941176 | 0.99215686 | 0.99215686 |
| 0.58823529 | 0.10588235 | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.0627451  | 0.36470588 | 0.98823529 | 0.99215686 | 0.73333333 |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.97647059 | 0.99215686 | 0.97647059 | 0.25098039 | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.18039216 | 0.50980392 | 0.71764706 | 0.99215686 |
| 0.99215686 | 0.81176471 | 0.00784314 | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.15294118 | 0.58039216 |
| 0.89803922 | 0.99215686 | 0.99215686 | 0.99215686 | 0.98039216 | 0.71372549 |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.09411765 | 0.44705882 | 0.86666667 | 0.99215686 | 0.99215686 | 0.99215686 |
| 0.99215686 | 0.78823529 | 0.30588235 | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.09019608 | 0.25882353 | 0.83529412 | 0.99215686 |
| 0.99215686 | 0.99215686 | 0.99215686 | 0.77647059 | 0.31764706 | 0.00784314 |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.07058824 | 0.67058824 |
| 0.85882353 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 | 0.76470588 |
| 0.31372549 | 0.03529412 | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.21568627 | 0.6745098  | 0.88627451 | 0.99215686 | 0.99215686 | 0.99215686 |
| 0.99215686 | 0.95686275 | 0.52156863 | 0.04313725 | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.         | 0.         |
| 0.         | 0.         | 0.         | 0.         | 0.53333333 | 0.99215686 |
| 0.99215686 | 0.99215686 | 0.83137255 | 0.52941176 | 0.51764706 | 0.0627451  |

[illegible]

## 2.2 Converting to One-Hot encoding

```
In [0]: # Here in this dataset, class labels for each image are numbers (0,1,2,3,...),
        # so I want to convert them into one-hot encoded vectors

        print("Class label of first image :", y_train[0])

        # Lets convert this into a 10 dimensional vector
        # ex: consider an image with class label of 5 and convert it into one-hot encoded vector of 0's and 1's - 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
        # this conversion is needed for MLPs

        Y_train = np_utils.to_categorical(y_train, 10)    # using np.utils we can convert numbers into one-hot encoding
        Y_test = np_utils.to_categorical(y_test, 10)

        print("After converting the output into a vector : ",Y_train[0])

        Class label of first image : 5
        After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```



```
In [0]: # this function is used draw Categorical Crossentropy Loss vs No. of epochs plot
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
import time

def plt_dynamic(x, vy, ty):
    plt.figure(figsize=(10,6))
    plt.plot(x, ty, 'b', label="Train Loss")
    plt.plot(x, vy, 'r', label="Validation/Test Loss")
    plt.title('\nCategorical Crossentropy Loss VS Epochs')
    plt.xlabel('Epochs')
    plt.ylabel('Categorical Crossentropy Loss-Train and Test loss')
    plt.legend()
    plt.grid()
    plt.show()
```

## 3. Softmax Classifier with 2 hidden layers

### 3.1 Model without dropout and Batch Normalization

```
In [0]: from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.initializers import he_normal

# some model parameters

output_dim = 10
input_dim = X_train.shape[1] # size of X
batch_size = 128
nb_epoch = 20    # run 20times
```

### 3.2 Building the model

```

In [0]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

# start building a model
# Initializing the sequential model as the layers are in sequential

model = Sequential()

# Adding first hidden layer
model.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))

# Adding second hidden layer
model.add(Dense(52, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding output layer
model.add(Dense(output_dim, activation='softmax'))

# Printing model summary
model.summary()

```

| Layer (type)              | Output Shape | Param # |
|---------------------------|--------------|---------|
| =====                     | =====        | =====   |
| dense_4 (Dense)           | (None, 364)  | 285740  |
| =====                     | =====        | =====   |
| dense_5 (Dense)           | (None, 52)   | 18980   |
| =====                     | =====        | =====   |
| dense_6 (Dense)           | (None, 10)   | 530     |
| =====                     | =====        | =====   |
| Total params: 305,250     |              |         |
| Trainable params: 305,250 |              |         |
| Non-trainable params: 0   |              |         |
| =====                     | =====        | =====   |

```
In [0]: # Compiling the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the model
# To fit the model, give input data, batch_size, number of epochs and validation/test data

history = model.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data = (X_test, Y_test))
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.cast instead.
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 100us/step - loss: 0.2683 - acc: 0.9218 - val_loss: 0.1305 - val_acc: 0.9612
Epoch 2/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1036 - acc: 0.9692 - val_loss: 0.0918 - val_acc: 0.9712
Epoch 3/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0662 - acc: 0.9802 - val_loss: 0.0769 - val_acc: 0.9753
Epoch 4/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0477 - acc: 0.9854 - val_loss: 0.0676 - val_acc: 0.9787
Epoch 5/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0340 - acc: 0.9898 - val_loss: 0.0809 - val_acc: 0.9751
Epoch 6/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0266 - acc: 0.9917 - val_loss: 0.0704 - val_acc: 0.9795
Epoch 7/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0197 - acc: 0.9946 - val_loss: 0.0673 - val_acc: 0.9802
Epoch 8/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0167 - acc: 0.9948 - val_loss: 0.0740 - val_acc: 0.9788
Epoch 9/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0161 - acc: 0.9948 - val_loss: 0.0659 - val_acc: 0.9820
Epoch 10/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0118 - acc: 0.9960 - val_loss: 0.0756 - val_acc: 0.9806
Epoch 11/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0112 - acc: 0.9964 - val_loss: 0.0737 - val_acc: 0.9817
Epoch 12/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0085 - acc: 0.9975 - val_loss: 0.0819 - val_acc: 0.9796
Epoch 13/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0070 - acc: 0.9980 - val_loss: 0.0904 - val_acc: 0.9778
Epoch 14/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0135 - acc: 0.9957 - val_loss: 0.0848 - val_acc: 0.9810
Epoch 15/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0065 - acc: 0.9980 - val_loss: 0.0906 - val_acc: 0.9791
Epoch 16/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0085 - acc: 0.9970 - val_loss: 0.0921 - val_acc: 0.9782
Epoch 17/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0084 - acc: 0.9973 - val_loss: 0.0968 - val_acc: 0.9791
```

```
Epoch 18/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0054 -
acc: 0.9982 - val_loss: 0.0825 - val_acc: 0.9827
Epoch 19/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0076 -
acc: 0.9975 - val_loss: 0.0943 - val_acc: 0.9803
Epoch 20/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0080 -
acc: 0.9974 - val_loss: 0.0915 - val_acc: 0.9820
```

### 3.3 Plotting the model values

```

In [0]: # Evaluating the model on test data
score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

# Plotting the results
# List of epoch numbers
x = list(range(1, nb_epoch+1))


# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
vy = history.history['val_loss']

# Training loss
ty = history.history['loss']
# calling the dynamic function to draw the plot
plt_dynamic(x, vy, ty)

```

Test score: 0.09153629648885012

Test accuracy: 0.982



### Observations:-

1. In this plot, as we are running more epochs the train and test loss are reducing, but test loss started increasing after epoch 9 and train loss/error is reducing further. This leads to model overfitting.
2. Actually test loss should decrease, but in above plot it's increasing and both train and test loss are diverging which leads to overfitting of model.
3. To avoid this, we can add regularizations like dropouts.

### 3.4 Plotting Violin plots of hidden and output layers to see weights distribution

This is like sanity check. This is needed to make sure that weights are not too large or too small.


```
In [0]: w_after = model.get_weights()

# weights of hidden Layer 1
h1_w = w_after[0].flatten().reshape(-1,1)
# weights of hidden Layer 2
h2_w = w_after[2].flatten().reshape(-1,1)
# weights of output Layer
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained\n")
# Hidden Layer 1
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

# Hidden Layer 2
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

# Output Layer
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



**Observation:-**

1. For the hidden layers, I got good gaussian structures and weights are nicely distributed . For output layer it is like combination of 2 gaussian structures .
2. All the weights are centered around 0 and have reasonable variance and all are working well.

## 3.5 Model with dropout and Batch Normalization



```

In [0]: from keras.layers.normalization import BatchNormalization
        from keras.layers import Dropout
        %matplotlib inline
        import warnings
        warnings.filterwarnings("ignore")

        model_drop = Sequential()

        # First hidden layer
        model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_
            initializer=he_normal(seed=None)))
        # Adding Batch Normalization
        model_drop.add(BatchNormalization())
        # Adding Dropout
        model_drop.add(Dropout(0.5))

        # Second hidden layer
        model_drop.add(Dense(52, activation='relu', kernel_initializer=he_normal(seed=
            None)))
        # Adding Batch Normalization
        model_drop.add(BatchNormalization())
        # Adding Dropout
        model_drop.add(Dropout(0.5))

        # Output Layer
        model_drop.add(Dense(output_dim, activation='softmax'))

        model_drop.summary()

```

| Layer (type)                                | Output Shape | Param # |
|---|--------------|---------|
| dense_10 (Dense)                            | (None, 364)  | 285740  |
| batch_normalization_3 (Batch Normalization) | (None, 364)  | 1456    |
| dropout_3 (Dropout)                         | (None, 364)  | 0       |
| dense_11 (Dense)                            | (None, 52)   | 18980   |
| batch_normalization_4 (Batch Normalization) | (None, 52)   | 208     |
| dropout_4 (Dropout)                         | (None, 52)   | 0       |
| dense_12 (Dense)                            | (None, 10)   | 530     |
| Total params: 306,914                       |              |         |
| Trainable params: 306,082                   |              |         |
| Non-trainable params: 832                   |              |         |

```
In [0]: # Compiling the model
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=
['accuracy'])

# Fitting the model

history = model_drop.fit(X_train, Y_train, batch_size = batch_size, epochs = n
b_epoch, verbose=1, validation_data = (X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 119us/step - loss: 0.5380 -  
acc: 0.8405 - val\_loss: 0.1708 - val\_acc: 0.9485

Epoch 2/20

60000/60000 [=====] - 6s 104us/step - loss: 0.2567 -  
acc: 0.9259 - val\_loss: 0.1255 - val\_acc: 0.9603

Epoch 3/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1998 -  
acc: 0.9431 - val\_loss: 0.1044 - val\_acc: 0.9678

Epoch 4/20

60000/60000 [=====] - 6s 105us/step - loss: 0.1731 -  
acc: 0.9509 - val\_loss: 0.0930 - val\_acc: 0.9717

Epoch 5/20

60000/60000 [=====] - 6s 108us/step - loss: 0.1526 -  
acc: 0.9566 - val\_loss: 0.0822 - val\_acc: 0.9737

Epoch 6/20

60000/60000 [=====] - 6s 106us/step - loss: 0.1421 -  
acc: 0.9596 - val\_loss: 0.0838 - val\_acc: 0.9728

Epoch 7/20

60000/60000 [=====] - 6s 102us/step - loss: 0.1304 -  
acc: 0.9622 - val\_loss: 0.0783 - val\_acc: 0.9771

Epoch 8/20

60000/60000 [=====] - 6s 101us/step - loss: 0.1237 -  
acc: 0.9643 - val\_loss: 0.0726 - val\_acc: 0.9787

Epoch 9/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1119 -  
acc: 0.9673 - val\_loss: 0.0741 - val\_acc: 0.9793

Epoch 10/20

60000/60000 [=====] - 6s 103us/step - loss: 0.1073 -  
acc: 0.9679 - val\_loss: 0.0728 - val\_acc: 0.9804

Epoch 11/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1046 -  
acc: 0.9693 - val\_loss: 0.0637 - val\_acc: 0.9809

Epoch 12/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0956 -  
acc: 0.9715 - val\_loss: 0.0711 - val\_acc: 0.9797

Epoch 13/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0924 -  
acc: 0.9732 - val\_loss: 0.0671 - val\_acc: 0.9805

Epoch 14/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0918 -  
acc: 0.9728 - val\_loss: 0.0699 - val\_acc: 0.9800

Epoch 15/20

60000/60000 [=====] - 6s 108us/step - loss: 0.0883 -  
acc: 0.9738 - val\_loss: 0.0667 - val\_acc: 0.9812

Epoch 16/20

60000/60000 [=====] - 6s 108us/step - loss: 0.0851 -  
acc: 0.9752 - val\_loss: 0.0667 - val\_acc: 0.9807

Epoch 17/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0802 -  
acc: 0.9756 - val\_loss: 0.0624 - val\_acc: 0.9816

Epoch 18/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0817 -  
acc: 0.9755 - val\_loss: 0.0622 - val\_acc: 0.9821

Epoch 19/20

60000/60000 [=====] - 7s 111us/step - loss: 0.0746 -

acc: 0.9779 - val\_loss: 0.0629 - val\_acc: 0.9823  
 Epoch 20/20  
 60000/60000 [=====] - 7s 110us/step - loss: 0.0776 -  
 acc: 0.9774 - val\_loss: 0.0625 - val\_acc: 0.9818

### 3.6 Plotting the model values

```
In [0]: # Evaluating the model on test data
score_drop = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_drop[0])
print('Test accuracy:', score_drop[1])


# Plotting the results
# list of epoch numbers
x = list(range(1, nb_epoch+1))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_Loss : validation loss
vy = history.history['val_loss']

# Training Loss
ty = history.history['loss']
# calling the dynamic function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.06253433417174965

Test accuracy: 0.9818



**Observation:-**

1. After adding dropout, both Train and Test loss are significantly reducing and there is no divergence / increasing between them .

**3.7 Plotting Violin plots of hidden and output layers to see weights distribution**

```

In [0]: w_after = model_drop.get_weights()

# weights of hidden Layer 1
h1_w = w_after[0].flatten().reshape(-1,1)
# weights of hidden Layer 2
h2_w = w_after[2].flatten().reshape(-1,1)
# weights of output Layer
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained\n")

# Hidden Layer 1
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

# Hidden Layer 2
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

# Output Layer
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## 4. Softmax Classifier with 3 hidden layers

### 4.1 Model without dropout and Batch Normalization

```

In [0]: # start building a model with 3 hidden layers
# Initializing the sequential model as the layers are in sequential
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

model_1 = Sequential()

# Adding first hidden layer
model_1.add(Dense(420, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))

# Adding second hidden layer
model_1.add(Dense(252, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding third hidden layer
model_1.add(Dense(45, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding output layer
model_1.add(Dense(output_dim, activation='softmax'))

# Printing model summary
model_1.summary()

```

| Layer (type)              | Output Shape | Param # |
|---------------------------|--------------|---------|
| =====                     | =====        | =====   |
| dense_13 (Dense)          | (None, 420)  | 329700  |
| dense_14 (Dense)          | (None, 252)  | 106092  |
| dense_15 (Dense)          | (None, 45)   | 11385   |
| dense_16 (Dense)          | (None, 10)   | 460     |
| =====                     | =====        | =====   |
| Total params: 447,637     |              |         |
| Trainable params: 447,637 |              |         |
| Non-trainable params: 0   |              |         |
| =====                     |              |         |

```
In [0]: # Compiling the model
model_1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the model
# To fit the model, give input data, batch_size, number of epochs and validation/test data

history = model_1.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data = (X_test, Y_test))
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 127us/step - loss: 0.2451 -  
acc: 0.9260 - val\_loss: 0.1085 - val\_acc: 0.9651

Epoch 2/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0893 -  
acc: 0.9726 - val\_loss: 0.0978 - val\_acc: 0.9687

Epoch 3/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0562 -  
acc: 0.9829 - val\_loss: 0.0833 - val\_acc: 0.9749

Epoch 4/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0410 -  
acc: 0.9870 - val\_loss: 0.0613 - val\_acc: 0.9819

Epoch 5/20

60000/60000 [=====] - 7s 117us/step - loss: 0.0295 -  
acc: 0.9902 - val\_loss: 0.0670 - val\_acc: 0.9798

Epoch 6/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0254 -  
acc: 0.9916 - val\_loss: 0.0775 - val\_acc: 0.9782

Epoch 7/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0219 -  
acc: 0.9926 - val\_loss: 0.0743 - val\_acc: 0.9794

Epoch 8/20

60000/60000 [=====] - 7s 116us/step - loss: 0.0193 -  
acc: 0.9938 - val\_loss: 0.0867 - val\_acc: 0.9770

Epoch 9/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0144 -  
acc: 0.9949 - val\_loss: 0.0889 - val\_acc: 0.9792

Epoch 10/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0146 -  
acc: 0.9948 - val\_loss: 0.0864 - val\_acc: 0.9780

Epoch 11/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0142 -  
acc: 0.9952 - val\_loss: 0.0838 - val\_acc: 0.9798

Epoch 12/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0126 -  
acc: 0.9959 - val\_loss: 0.0747 - val\_acc: 0.9818

Epoch 13/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0103 -  
acc: 0.9966 - val\_loss: 0.0833 - val\_acc: 0.9801

Epoch 14/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0106 -  
acc: 0.9960 - val\_loss: 0.0899 - val\_acc: 0.9806

Epoch 15/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0115 -  
acc: 0.9963 - val\_loss: 0.0983 - val\_acc: 0.9788

Epoch 16/20

60000/60000 [=====] - 7s 116us/step - loss: 0.0081 -  
acc: 0.9975 - val\_loss: 0.0873 - val\_acc: 0.9803

Epoch 17/20

60000/60000 [=====] - 7s 114us/step - loss: 0.0085 -  
acc: 0.9974 - val\_loss: 0.0980 - val\_acc: 0.9797

Epoch 18/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0089 -  
acc: 0.9971 - val\_loss: 0.1002 - val\_acc: 0.9792

Epoch 19/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0106 -

acc: 0.9966 - val\_loss: 0.0963 - val\_acc: 0.9818

Epoch 20/20

60000/60000 [=====] - 7s 116us/step - loss: 0.0088 -

acc: 0.9972 - val\_loss: 0.1031 - val\_acc: 0.9807

## 4.2 Plotting the model values

```

In [0]: # Evaluating the model on test data
score_1 = model_1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_1[0])
print('Test accuracy:', score_1[1])

# Plotting the results
# list of epoch numbers
x = list(range(1, nb_epoch+1))

# https://machinelearningmastery.com/display-deep-learning-model-training-history-in-keras/
# we will get val_loss and val_acc only when you pass the parameter validation_data
# for each key in history.history we will have a list of length equal to number of epochs. History records training metrics for each epoch.
# This includes the loss and the accuracy as well as the loss and accuracy for the validation dataset, if one is set.
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

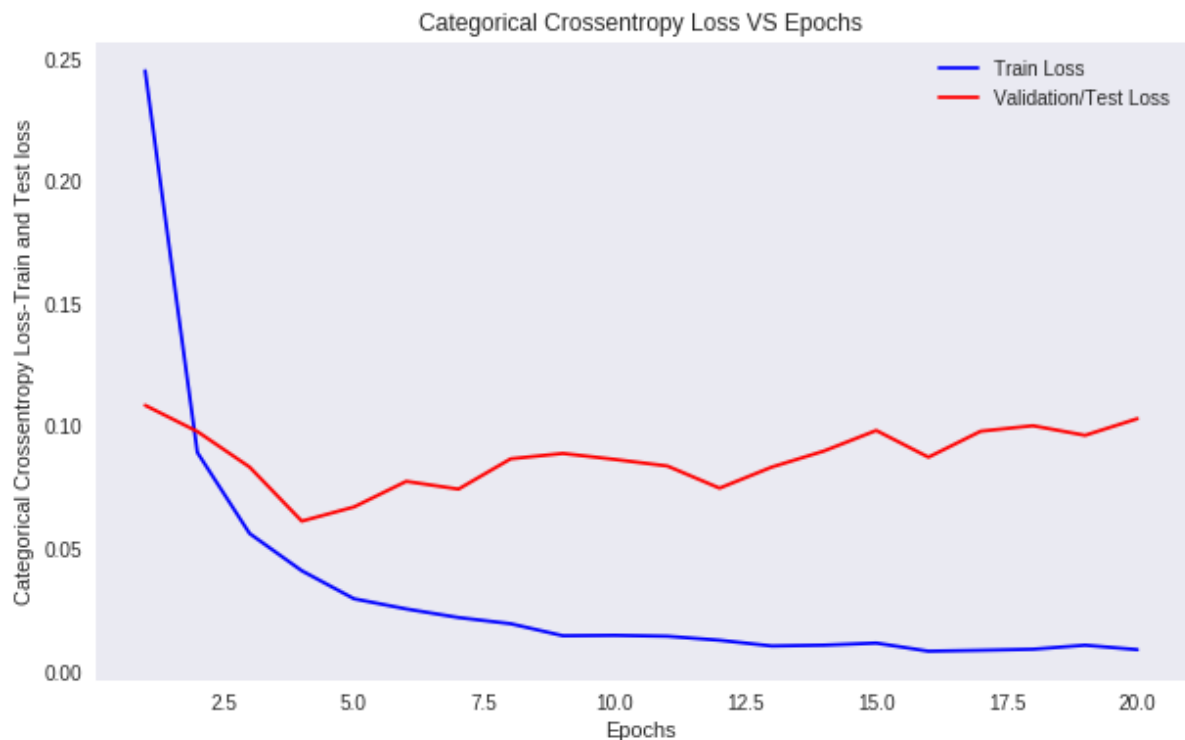
# val_loss : validation loss
vy = history.history['val_loss']
# Training loss
ty = history.history['loss']

# calling the dynamic function to draw the plot
plt_dynamic(x, vy, ty)

```

Test score: 0.10307921019351929

Test accuracy: 0.9807



**Observations:-**

1. In this plot, as we are running more epochs the train and test loss are reducing , but test loss started increasing after epoch 6 and train loss/error is reducing further. This leads to model overfitting
2. Actually test loss should decrease, but in above plot its increasing and both train and test loss are diverging which leads to overfitting of model.
3. To avoid this, we can add regularizations like dropouts.

**4.3 Plotting Violin plots of hidden and output layers to see weights distribution**

```

In [0]: w_after = model_1.get_weights()

# 3 hidden layers and output layer
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained\n")


# For hidden layer 1
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

# For hidden layer 2
plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

#For hidden layer 3
plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## 4.4 Model with dropout and Batch Normalization

```
In [0]: from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

model_1_drop = Sequential()

# First hidden layer
model_1_drop.add(Dense(420, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_1_drop.add(BatchNormalization())
# Adding Dropout
model_1_drop.add(Dropout(0.5))

# Second hidden layer
model_1_drop.add(Dense(252, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_1_drop.add(BatchNormalization())
# Adding Dropout
model_1_drop.add(Dropout(0.5))

# Third hidden layer
model_1_drop.add(Dense(45, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_1_drop.add(BatchNormalization())
# Adding Dropout
model_1_drop.add(Dropout(0.5))

# Output layer
model_1_drop.add(Dense(output_dim, activation='softmax'))

model_1_drop.summary()
```

| Layer (type)                                | Output Shape | Param # |
|---|--------------|---------|
| =====                                       | =====        | =====   |
| dense_17 (Dense)                            | (None, 420)  | 329700  |
| batch_normalization_5 (Batch Normalization) | (None, 420)  | 1680    |
| dropout_5 (Dropout)                         | (None, 420)  | 0       |
| dense_18 (Dense)                            | (None, 252)  | 106092  |
| batch_normalization_6 (Batch Normalization) | (None, 252)  | 1008    |
| dropout_6 (Dropout)                         | (None, 252)  | 0       |
| dense_19 (Dense)                            | (None, 45)   | 11385   |
| batch_normalization_7 (Batch Normalization) | (None, 45)   | 180     |
| dropout_7 (Dropout)                         | (None, 45)   | 0       |
| dense_20 (Dense)                            | (None, 10)   | 460     |
| =====                                       | =====        | =====   |
| Total params: 450,505                       |              |         |
| Trainable params: 449,071                   |              |         |
| Non-trainable params: 1,434                 |              |         |
| =====                                       | =====        | =====   |

```
In [0]: # Compiling the model
model_1_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the model

history = model_1_drop.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data = (X_test, Y_test))
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 181us/step - loss: 0.6888  
- acc: 0.7916 - val\_loss: 0.1892 - val\_acc: 0.9412

Epoch 2/20

60000/60000 [=====] - 10s 160us/step - loss: 0.2975  
- acc: 0.9179 - val\_loss: 0.1274 - val\_acc: 0.9596

Epoch 3/20

60000/60000 [=====] - 9s 158us/step - loss: 0.2302 -  
acc: 0.9372 - val\_loss: 0.1156 - val\_acc: 0.9645

Epoch 4/20

60000/60000 [=====] - 9s 158us/step - loss: 0.1955 -  
acc: 0.9456 - val\_loss: 0.0927 - val\_acc: 0.9728

Epoch 5/20

60000/60000 [=====] - 9s 157us/step - loss: 0.1660 -  
acc: 0.9546 - val\_loss: 0.0977 - val\_acc: 0.9714

Epoch 6/20

60000/60000 [=====] - 10s 158us/step - loss: 0.1559  
- acc: 0.9566 - val\_loss: 0.0861 - val\_acc: 0.9738

Epoch 7/20

60000/60000 [=====] - 10s 159us/step - loss: 0.1418  
- acc: 0.9600 - val\_loss: 0.0807 - val\_acc: 0.9765

Epoch 8/20

60000/60000 [=====] - 10s 159us/step - loss: 0.1333  
- acc: 0.9637 - val\_loss: 0.0731 - val\_acc: 0.9785

Epoch 9/20

60000/60000 [=====] - 9s 156us/step - loss: 0.1274 -  
acc: 0.9652 - val\_loss: 0.0800 - val\_acc: 0.9779

Epoch 10/20

60000/60000 [=====] - 9s 156us/step - loss: 0.1168 -  
acc: 0.9673 - val\_loss: 0.0784 - val\_acc: 0.9777

Epoch 11/20

60000/60000 [=====] - 9s 156us/step - loss: 0.1105 -  
acc: 0.9688 - val\_loss: 0.0695 - val\_acc: 0.9787

Epoch 12/20

60000/60000 [=====] - 9s 156us/step - loss: 0.1045 -  
acc: 0.9703 - val\_loss: 0.0732 - val\_acc: 0.9795

Epoch 13/20

60000/60000 [=====] - 9s 157us/step - loss: 0.1041 -  
acc: 0.9709 - val\_loss: 0.0678 - val\_acc: 0.9806

Epoch 14/20

60000/60000 [=====] - 10s 159us/step - loss: 0.0973  
- acc: 0.9723 - val\_loss: 0.0706 - val\_acc: 0.9805

Epoch 15/20

60000/60000 [=====] - 10s 159us/step - loss: 0.0959  
- acc: 0.9742 - val\_loss: 0.0664 - val\_acc: 0.9816

Epoch 16/20

60000/60000 [=====] - 10s 159us/step - loss: 0.0902  
- acc: 0.9739 - val\_loss: 0.0658 - val\_acc: 0.9822

Epoch 17/20

60000/60000 [=====] - 9s 158us/step - loss: 0.0853 -  
acc: 0.9755 - val\_loss: 0.0657 - val\_acc: 0.9832

Epoch 18/20

60000/60000 [=====] - 9s 157us/step - loss: 0.0838 -  
acc: 0.9759 - val\_loss: 0.0655 - val\_acc: 0.9819

Epoch 19/20

60000/60000 [=====] - 9s 156us/step - loss: 0.0827 -

acc: 0.9771 - val\_loss: 0.0608 - val\_acc: 0.9839  
 Epoch 20/20  
 60000/60000 [=====] - 9s 153us/step - loss: 0.0801 -  
 acc: 0.9774 - val\_loss: 0.0603 - val\_acc: 0.9834

## 4.5 Plotting the model values

```
In [0]: # Evaluating the model on test data
score_1_drop = model_1_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_1_drop[0])
print('Test accuracy:', score_1_drop[1])


# Plotting the results
# list of epoch numbers
x = list(range(1, nb_epoch+1))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_Loss : validation loss
vy = history.history['val_loss']

# Training loss
ty = history.history['loss']
# calling the dynamic function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.0603378255185904

Test accuracy: 0.9834



**Observation:-**

After adding dropout, both Train and Test loss are significantly reducing /converging and there is no divergence / increasing between them .

**4.6 Plotting Violin plots of hidden and output layers to see weights distribution**

```

In [0]: w_after = model_1_drop.get_weights()

# weights of hidden layer 1
h1_w = w_after[0].flatten().reshape(-1,1)
# weights of hidden Layer 2
h2_w = w_after[2].flatten().reshape(-1,1)
# weights of hidden Layer 3
h3_w = w_after[4].flatten().reshape(-1,1)
# weights of output Layer
out_w = w_after[6].flatten().reshape(-1,1)

# Plotting the violin plots
fig = plt.figure()
plt.title("Weight matrices after model trained\n")


# Hidden Layer 1
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

# Hidden Layer 2
plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

# Hidden Layer 3
plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

# Output Layer
plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## **5. Softmax Classifier with 5 hidden layers**

### **5.1 Model without dropout and Batch Normalization**

```

In [0]: # start building a model with 5 hidden layers
# Initializing the sequential model as the layers are in sequential
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

model_2 = Sequential()

# Adding first hidden layer
model_2.add(Dense(532, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))

# Adding second hidden layer
model_2.add(Dense(354, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding third hidden layer
model_2.add(Dense(165, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding fourth hidden layer
model_2.add(Dense(83, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding fifth hidden layer
model_2.add(Dense(34, activation='relu', kernel_initializer=he_normal(seed=None)))

# Adding output layer
model_2.add(Dense(output_dim, activation='softmax'))

# Printing model summary
model_2.summary()

```

| Layer (type)              | Output Shape | Param # |
|---------------------------|--------------|---------|
| =====                     |              |         |
| dense_21 (Dense)          | (None, 532)  | 417620  |
| dense_22 (Dense)          | (None, 354)  | 188682  |
| dense_23 (Dense)          | (None, 165)  | 58575   |
| dense_24 (Dense)          | (None, 83)   | 13778   |
| dense_25 (Dense)          | (None, 34)   | 2856    |
| dense_26 (Dense)          | (None, 10)   | 350     |
| =====                     |              |         |
| Total params: 681,861     |              |         |
| Trainable params: 681,861 |              |         |
| Non-trainable params: 0   |              |         |
| =====                     |              |         |

```
In [0]: # Compiling the model
model_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the model
# To fit the model, give input data, batch_size, number of epochs and validation/test data

history = model_2.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data = (X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 192us/step - loss: 0.2578  
- acc: 0.9197 - val\_loss: 0.1270 - val\_acc: 0.9578

Epoch 2/20

60000/60000 [=====] - 10s 170us/step - loss: 0.0911  
- acc: 0.9718 - val\_loss: 0.0933 - val\_acc: 0.9718

Epoch 3/20

60000/60000 [=====] - 10s 173us/step - loss: 0.0614  
- acc: 0.9804 - val\_loss: 0.0871 - val\_acc: 0.9730

Epoch 4/20

60000/60000 [=====] - 10s 171us/step - loss: 0.0464  
- acc: 0.9846 - val\_loss: 0.0854 - val\_acc: 0.9757

Epoch 5/20

60000/60000 [=====] - 10s 171us/step - loss: 0.0380  
- acc: 0.9879 - val\_loss: 0.0770 - val\_acc: 0.9785

Epoch 6/20

60000/60000 [=====] - 10s 171us/step - loss: 0.0292  
- acc: 0.9905 - val\_loss: 0.0789 - val\_acc: 0.9766

Epoch 7/20

60000/60000 [=====] - 10s 172us/step - loss: 0.0293  
- acc: 0.9906 - val\_loss: 0.0700 - val\_acc: 0.9821

Epoch 8/20

60000/60000 [=====] - 11s 176us/step - loss: 0.0253  
- acc: 0.9919 - val\_loss: 0.0860 - val\_acc: 0.9797

Epoch 9/20

60000/60000 [=====] - 10s 174us/step - loss: 0.0221  
- acc: 0.9925 - val\_loss: 0.0784 - val\_acc: 0.9811

Epoch 10/20

60000/60000 [=====] - 10s 172us/step - loss: 0.0181  
- acc: 0.9938 - val\_loss: 0.0826 - val\_acc: 0.9796

Epoch 11/20

60000/60000 [=====] - 10s 173us/step - loss: 0.0163  
- acc: 0.9950 - val\_loss: 0.0972 - val\_acc: 0.9775

Epoch 12/20

60000/60000 [=====] - 10s 172us/step - loss: 0.0161  
- acc: 0.9948 - val\_loss: 0.0794 - val\_acc: 0.9800

Epoch 13/20

60000/60000 [=====] - 10s 173us/step - loss: 0.0148  
- acc: 0.9952 - val\_loss: 0.0783 - val\_acc: 0.9834

Epoch 14/20

60000/60000 [=====] - 10s 171us/step - loss: 0.0163  
- acc: 0.9950 - val\_loss: 0.0939 - val\_acc: 0.9808

Epoch 15/20

60000/60000 [=====] - 10s 171us/step - loss: 0.0121  
- acc: 0.9964 - val\_loss: 0.0874 - val\_acc: 0.9799

Epoch 16/20

60000/60000 [=====] - 10s 171us/step - loss: 0.0130  
- acc: 0.9962 - val\_loss: 0.0859 - val\_acc: 0.9809

Epoch 17/20

60000/60000 [=====] - 10s 169us/step - loss: 0.0114  
- acc: 0.9969 - val\_loss: 0.0865 - val\_acc: 0.9824

Epoch 18/20

60000/60000 [=====] - 10s 170us/step - loss: 0.0129  
- acc: 0.9962 - val\_loss: 0.0752 - val\_acc: 0.9836

Epoch 19/20

60000/60000 [=====] - 10s 170us/step - loss: 0.0106



```
- acc: 0.9968 - val_loss: 0.0830 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 10s 173us/step - loss: 0.0097
- acc: 0.9970 - val_loss: 0.0877 - val_acc: 0.9824
```

## 5.2 Plotting the model values

```
In [0]: # Evaluating the model on test data
score_2 = model_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_2[0])
print('Test accuracy:', score_2[1])

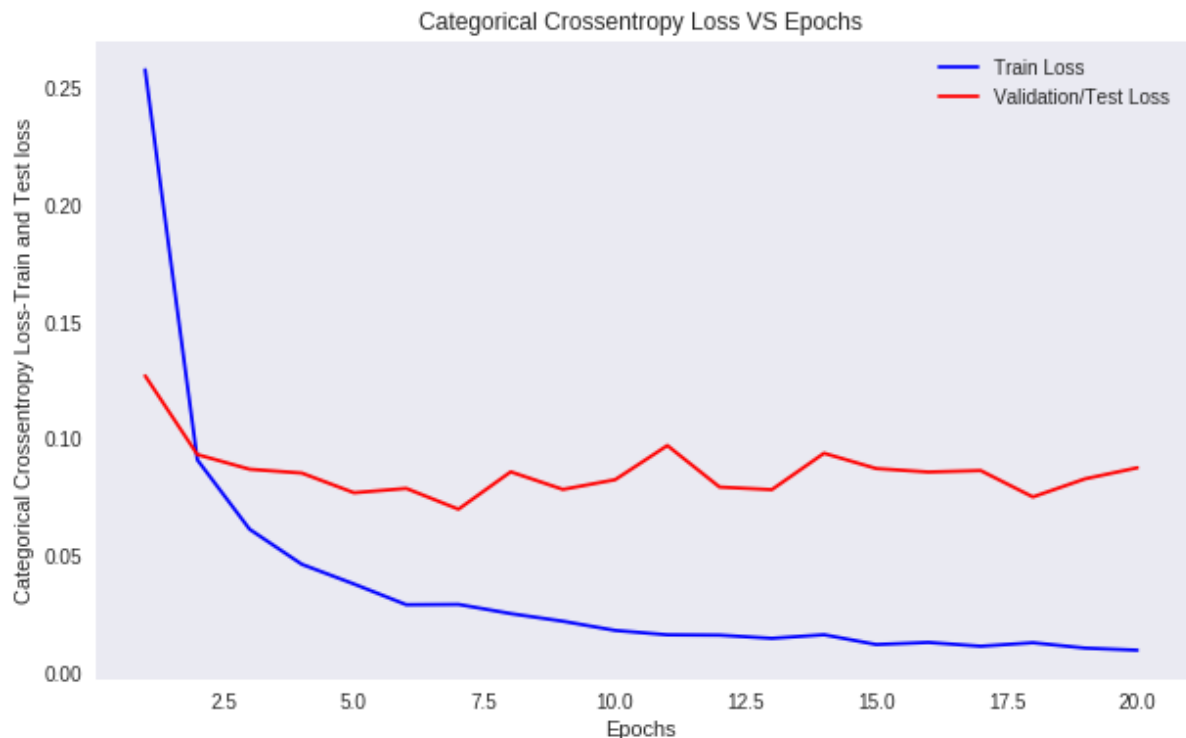
# Plotting the results
# list of epoch numbers
x = list(range(1, nb_epoch+1))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_Loss : validation loss
vy = history.history['val_loss']

# Training Loss
ty = history.history['loss']
# calling the dynamic function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.08769175834859744

Test accuracy: 0.9824



### **5.3 Plotting Violin plots of hidden and output layers to see weights distribution**

```
In [0]: w_after = model_2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)

out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained\n")

# Hidden Layer 1
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

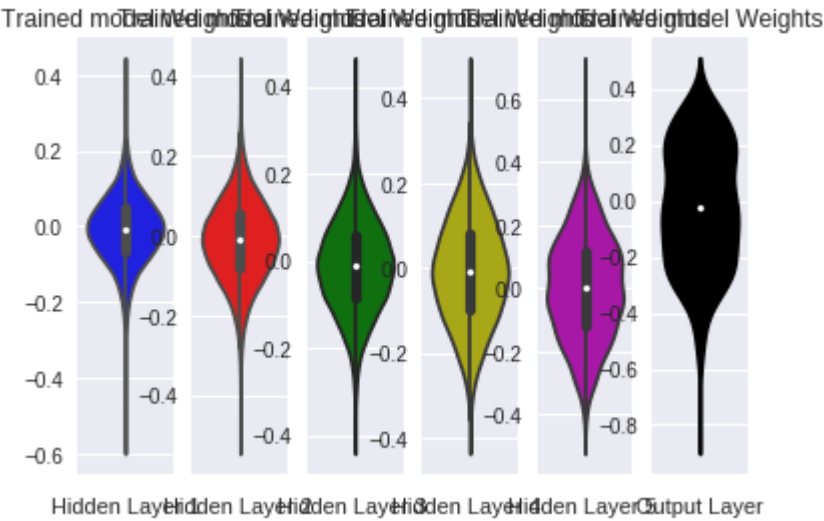
# Hidden Layer 2
plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

# Hidden Layer 3
plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

# Hidden Layer 4
plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='y')
plt.xlabel('Hidden Layer 4 ')

# Hidden Layer 5
plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='m')
plt.xlabel('Hidden Layer 5 ')

# Output Layer
plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='k')
plt.xlabel('Output Layer')
plt.show()
```



### 5.4 Model with dropout and Batch Normalization

```
In [0]: from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

model_2_drop = Sequential()

# First hidden layer
model_2_drop.add(Dense(532, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_2_drop.add(BatchNormalization())
# Adding Dropout
model_2_drop.add(Dropout(0.5))

# Second hidden layer
model_2_drop.add(Dense(354, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_2_drop.add(BatchNormalization())
# Adding Dropout
model_2_drop.add(Dropout(0.5))

# Third hidden layer
model_2_drop.add(Dense(165, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_2_drop.add(BatchNormalization())
# Adding Dropout
model_2_drop.add(Dropout(0.5))

# Fouth hidden layer
model_2_drop.add(Dense(83, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_2_drop.add(BatchNormalization())
# Adding Dropout
model_2_drop.add(Dropout(0.5))

# Fifth hidden layer
model_2_drop.add(Dense(34, activation='relu', kernel_initializer=he_normal(seed=None)))
# Adding Batch Normalization
model_2_drop.add(BatchNormalization())
# Adding Dropout
model_2_drop.add(Dropout(0.5))

# Output layer
model_2_drop.add(Dense(output_dim, activation='softmax'))

model_2_drop.summary()
```

| Layer (type)                                 | Output Shape | Param # |
|--|--------------|---------|
| dense_27 (Dense)                             | (None, 532)  | 417620  |
| batch_normalization_8 (Batch Normalization)  | (None, 532)  | 2128    |
| dropout_8 (Dropout)                          | (None, 532)  | 0       |
| dense_28 (Dense)                             | (None, 354)  | 188682  |
| batch_normalization_9 (Batch Normalization)  | (None, 354)  | 1416    |
| dropout_9 (Dropout)                          | (None, 354)  | 0       |
| dense_29 (Dense)                             | (None, 165)  | 58575   |
| batch_normalization_10 (Batch Normalization) | (None, 165)  | 660     |
| dropout_10 (Dropout)                         | (None, 165)  | 0       |
| dense_30 (Dense)                             | (None, 83)   | 13778   |
| batch_normalization_11 (Batch Normalization) | (None, 83)   | 332     |
| dropout_11 (Dropout)                         | (None, 83)   | 0       |
| dense_31 (Dense)                             | (None, 34)   | 2856    |
| batch_normalization_12 (Batch Normalization) | (None, 34)   | 136     |
| dropout_12 (Dropout)                         | (None, 34)   | 0       |
| dense_32 (Dense)                             | (None, 10)   | 350     |
| Total params: 686,533                        |              |         |
| Trainable params: 684,197                    |              |         |
| Non-trainable params: 2,336                  |              |         |

```
In [0]: # Compiling the model
model_2_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the model

history = model_2_drop.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=1, validation_data = (X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 276us/step - loss: 1.3228  
- acc: 0.5787 - val\_loss: 0.2892 - val\_acc: 0.9192

Epoch 2/20

60000/60000 [=====] - 14s 237us/step - loss: 0.5058  
- acc: 0.8595 - val\_loss: 0.1904 - val\_acc: 0.9490

Epoch 3/20

60000/60000 [=====] - 14s 237us/step - loss: 0.3545  
- acc: 0.9100 - val\_loss: 0.1504 - val\_acc: 0.9596

Epoch 4/20

60000/60000 [=====] - 14s 235us/step - loss: 0.2935  
- acc: 0.9292 - val\_loss: 0.1263 - val\_acc: 0.9661

Epoch 5/20

60000/60000 [=====] - 14s 235us/step - loss: 0.2576  
- acc: 0.9379 - val\_loss: 0.1281 - val\_acc: 0.9682

Epoch 6/20

60000/60000 [=====] - 14s 234us/step - loss: 0.2308  
- acc: 0.9458 - val\_loss: 0.1127 - val\_acc: 0.9721

Epoch 7/20

60000/60000 [=====] - 14s 236us/step - loss: 0.2067  
- acc: 0.9514 - val\_loss: 0.1093 - val\_acc: 0.9715

Epoch 8/20

60000/60000 [=====] - 14s 236us/step - loss: 0.1912  
- acc: 0.9553 - val\_loss: 0.1021 - val\_acc: 0.9747

Epoch 9/20

60000/60000 [=====] - 14s 235us/step - loss: 0.1823  
- acc: 0.9578 - val\_loss: 0.0959 - val\_acc: 0.9762

Epoch 10/20

60000/60000 [=====] - 14s 237us/step - loss: 0.1711  
- acc: 0.9598 - val\_loss: 0.0918 - val\_acc: 0.9762

Epoch 11/20

60000/60000 [=====] - 14s 237us/step - loss: 0.1628  
- acc: 0.9613 - val\_loss: 0.0918 - val\_acc: 0.9775

Epoch 12/20

60000/60000 [=====] - 15s 242us/step - loss: 0.1514  
- acc: 0.9648 - val\_loss: 0.0931 - val\_acc: 0.9775

Epoch 13/20

60000/60000 [=====] - 15s 250us/step - loss: 0.1484  
- acc: 0.9659 - val\_loss: 0.0922 - val\_acc: 0.9784

Epoch 14/20

60000/60000 [=====] - 14s 235us/step - loss: 0.1397  
- acc: 0.9673 - val\_loss: 0.0853 - val\_acc: 0.9799

Epoch 15/20

60000/60000 [=====] - 14s 237us/step - loss: 0.1390  
- acc: 0.9675 - val\_loss: 0.0811 - val\_acc: 0.9806

Epoch 16/20

60000/60000 [=====] - 14s 240us/step - loss: 0.1280  
- acc: 0.9700 - val\_loss: 0.0786 - val\_acc: 0.9809

Epoch 17/20

60000/60000 [=====] - 14s 239us/step - loss: 0.1215  
- acc: 0.9716 - val\_loss: 0.0842 - val\_acc: 0.9800

Epoch 18/20

60000/60000 [=====] - 14s 236us/step - loss: 0.1196  
- acc: 0.9721 - val\_loss: 0.0788 - val\_acc: 0.9821

Epoch 19/20

60000/60000 [=====] - 14s 237us/step - loss: 0.1168



- acc: 0.9725 - val\_loss: 0.0754 - val\_acc: 0.9808

Epoch 20/20

60000/60000 [=====] - 14s 235us/step - loss: 0.1139

- acc: 0.9728 - val\_loss: 0.0776 - val\_acc: 0.9815

## 5.5 Plotting the model values

```
In [0]: # Evaluating the model on test data
score_2_drop = model_2_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score_2_drop[0])
print('Test accuracy:', score_2_drop[1])

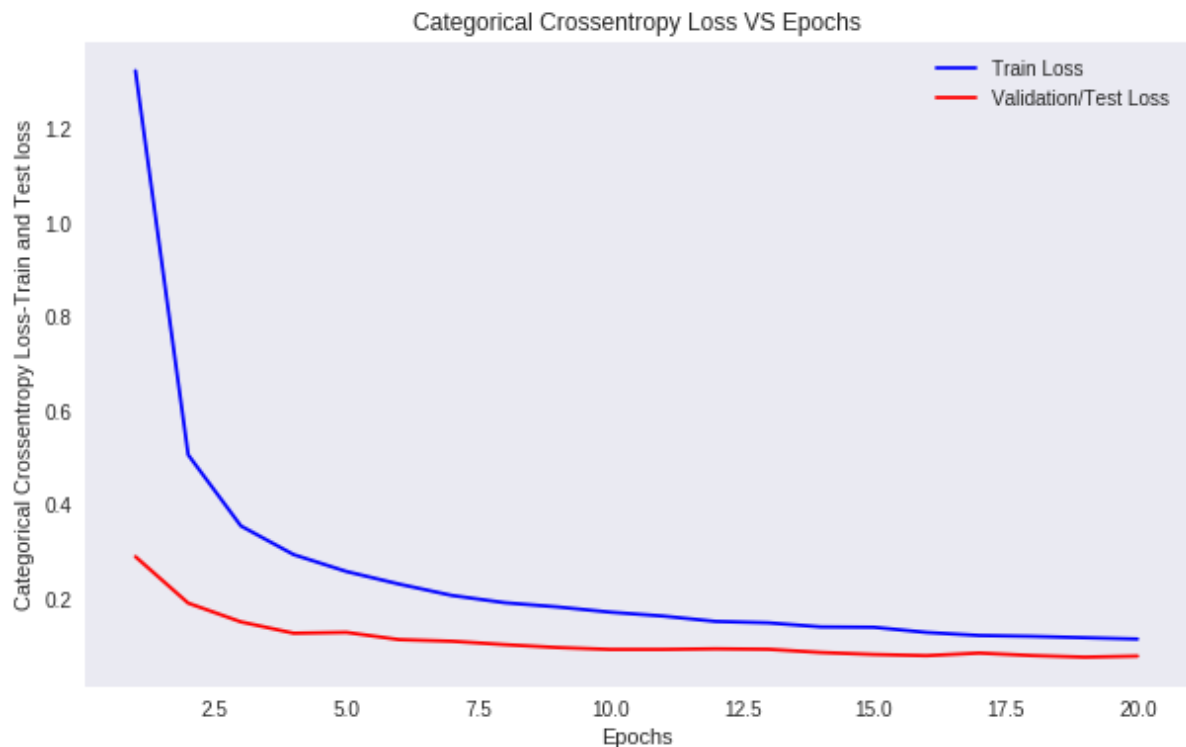
# Plotting the results
# list of epoch numbers
x = list(range(1, nb_epoch+1))

# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_Loss : validation loss
vy = history.history['val_loss']

# Training loss
ty = history.history['loss']
# calling the dynamic function to draw the plot
plt_dynamic(x, vy, ty)
```

Test score: 0.07762294698476326

Test accuracy: 0.9815



## 5.6 Plotting Violin plots of hidden and output layers to see weights distribution

```
In [0]: w_after = model_2_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)

out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained\n")

# Hidden Layer 1
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

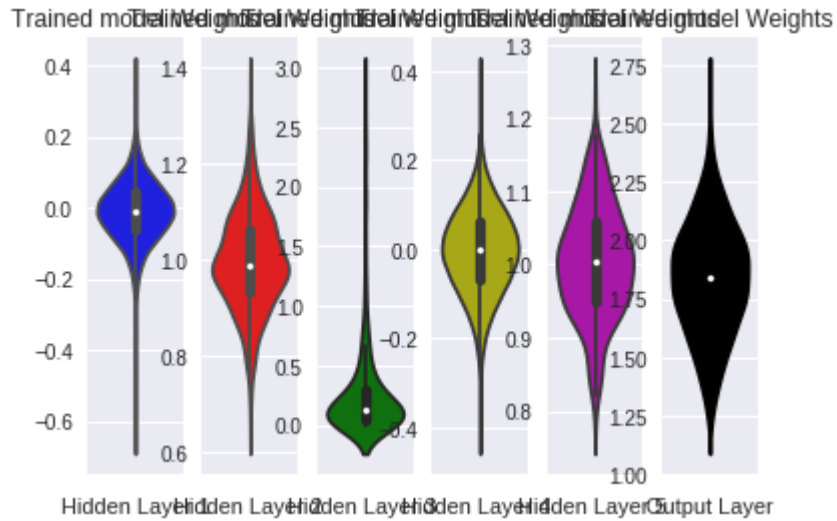
# Hidden Layer 2
plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

# Hidden Layer 3
plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

# Hidden Layer 4
plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='y')
plt.xlabel('Hidden Layer 4 ')

# Hidden Layer 5
plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='m')
plt.xlabel('Hidden Layer 5 ')

# Output Layer
plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='k')
plt.xlabel('Output Layer')
plt.show()
```



## 6. Models Summarization

```
In [0]: from pandas import DataFrame
MLP = {'MLP Layers':['2-hidden layer','2-hidden layer','3-hidden layer','3-hidden layer','5-hidden layer','5-hidden layer'],
      'Model':['Model without dropout and Batch Normalization','Model with dropout and Batch Normalization','Model without dropout and Batch Normalization',
               'Model with dropout and Batch Normalization','Model without dropout and Batch Normalization'],
      'Activation':['ReLU','ReLU','ReLU','ReLU','ReLU','ReLU'],
      'Optimizer':['Adam','Adam','Adam','Adam','Adam','Adam'],
      'Kernel_initializer':['he_normal','he_normal','he_normal','he_normal','he_normal','he_normal'],
      'output layer':['softmax','softmax','softmax','softmax','softmax','softmax'],
      'Training accuracy':['0.99','0.97','0.99','0.97','0.99','0.97'],
      'Train loss':['0.008','0.07','0.008','0.08','0.009','0.11'],
      'Test accuracy':['0.98','0.98','0.98','0.98','0.98','0.98'],
      'Test loss':['0.091','0.06','0.10','0.06','0.08','0.07']}
```

```
In [0]: Final_conclusions = DataFrame(MLP)
Final_conclusions
```

```
Out[0]:
```

|          | Activation | Kernel_initializer | MLP<br>Layers         | Model   | Optimizer | Test<br>accuracy | Test<br>loss | Train<br>loss |
|----------|------------|--------------------|-----------------------|---|-----------|------------------|--------------|---------------|
| <b>0</b> | ReLU       | he_normal          | 2-<br>hidden<br>layer | Model<br>without<br>dropout and<br>Batch<br>Normalization | Adam      | 0.98             | 0.091        | 0.008         |
| <b>1</b> | ReLU       | he_normal          | 2-<br>hidden<br>layer | Model with<br>dropout and<br>Batch<br>Normalization       | Adam      | 0.98             | 0.06         | 0.07          |
| <b>2</b> | ReLU       | he_normal          | 3-<br>hidden<br>layer | Model<br>without<br>dropout and<br>Batch<br>Normalization | Adam      | 0.98             | 0.10         | 0.008         |
| <b>3</b> | ReLU       | he_normal          | 3-<br>hidden<br>layer | Model with<br>dropout and<br>Batch<br>Normalization       | Adam      | 0.98             | 0.06         | 0.08          |
| <b>4</b> | ReLU       | he_normal          | 5-<br>hidden<br>layer | Model<br>without<br>dropout and<br>Batch<br>Normalization | Adam      | 0.98             | 0.08         | 0.009         |
| <b>5</b> | ReLU       | he_normal          | 5-<br>hidden<br>layer | Model with<br>dropout and<br>Batch<br>Normalization       | Adam      | 0.98             | 0.07         | 0.11          |



## 7. Conclusions:-

From the above observations I can observed,

1. As I had used ReLU activation and Adam optimizer, all the accuracies are good.
2. Output layer is softmax layer
3. Models without dropout and normalization have less Test accuracy, and after adding dropout and normalization, model Test/validation accuracy improved, but there is no much difference in values.
4. With the add of dropout and normalization, the test loss also decreased.
5. In the Categorical Crossentropy Loss VS Epochs plot, train and test loss also converged or decreased gradually after adding dropout and normalization.
6. So by adding dropout and normalization, model worked well and gave good results.
7. I have used he\_normal as kernel initializer as it have inbuilt mean and std values, without need of defining them explicitly.